

Intensity Transformations and Neighborhood Filtering

GitHub Link :- <https://github.com/Manimohan05/Image-Processing-and-Machine-Vision--A01-Intensity-Transformations-and-Neighborhood-Filtering.git>

1. Question 01

In this question, we must implement an intensity transformation on the input image. Through the adjustment of pixel values using this transformation, our objective is to enhance the overall visual quality and appearance of the image.

Important part of the code

```
c = np.array([(100, 50), (150, 200)])
arr1 = np.linspace(0, c[0, 1], c[0, 1] + 1 - 0).astype('uint8')
arr2 = np.linspace(c[0, 0] + 1, 255, c[1, 1] - c[0, 0]).astype('uint8')
arr3 = np.linspace(c[1, 0] + 1, 255, 255 - c[1, 0]).astype('uint8')
transform = np.concatenate((arr1, arr2), axis=0).astype('uint8')
transform = np.concatenate((transform, arr3), axis=0).astype('uint8')
image_transformed = cv.LUT(img_orig, transform)
```




figure 1:- Intensity transformation



figure 2:- Original and Transformed

2. Question 02

Our task involves carrying out a comparable procedure to the one mentioned earlier, with the goal of highlighting specific brain tissue features in proton density images. More precisely, our objective is to improve the perceptibility of both white matter and gray matter separately. This requires developing personalized intensity adjustments specific to each type of tissue. To record this procedure, we will generate plots of intensity transformations that visually depict the modifications applied to enhance the visibility of white matter and gray matter within the image.

Important code for accentuating gray matter.

```
c = np.array([(0,50),(150,100),(180,255)])
arr1=np.linspace(0,c[1,1],c[0,1]+1).astype("uint8")
arr2 = np.linspace(c[2,0]+1,c[2,1],c[2,0]-c[0,1]).astype("uint8")
arr3 = np.linspace(c[0,1]-20,c[2,0]-c[0,1],255-c[2,0]).astype("uint8")
gray_matter_transformation = np.concatenate((arr1,arr2),axis=0).astype("uint8")
gray_matter_transformation = np.concatenate((gray_matter_transformation,arr3),axis=0).astype("uint8")
```




figure 3:- Intensity transformation(gray)

Important code for accentuating white matter.

```
c=np.array([(50,25),(50,90),(180,220),(255,255)])
```

```

arr1=np.linspace(0,c[0,1],c[0,0]+1).astype("uint8")
arr2 = np.linspace(c[0,1]+1,c[1,1],c[2,0]-
c[1,0]).astype("uint8")
arr3 = np.linspace(c[2,1]+1,255,255-
c[2,0]).astype("uint8")
white_matter_transformation =
np.concatenate((arr1,arr2),axis=0).astype("uint8")
white_matter_transformation=
np.concatenate((white_matter_transformation,arr3),axis=0)
.astype("uint8")

```




figure 4:- Intensity transformation(white)




figure 5:- Original and Transformed

3. Question 03

In this task, we are utilizing gamma correction on the L plane within the L^*a^*b color space. Gamma correction is a method employed to fine-tune the intensity levels of an image, and in this instance, it will be employed on the luminance channel (L) of the color space. We will be given a particular γ value that determines the degree of correction. Furthermore, we will create histograms to visually contrast the intensity distribution of the original L channel with that of the adjusted version. These histograms provide a means to assess the level of improvement accomplished through the gamma correction procedure.

Important part of code

```

# Load the original image
img_orig = cv.imread('images\highlights_and_shadows.jpg',
cv.IMREAD_COLOR)
# Convert the image to LAB color space
img_lab = cv.cvtColor(img_orig, cv.COLOR_BGR2LAB)
# Apply gamma correction to the L channel
gamma = 1.4
img_lab[:, :, 0] =
(((img_orig[:, :, 0]/255)**gamma)*255).astype(np.uint8)
# Convert the image back to BGR color space
img_gamma = cv.cvtColor(img_lab, cv.COLOR_LAB2BGR)

```




figure 6:- Original and L plane corrected Images




figure 7:- Histograms

4. Question 04

To enhance the vibrancy of a photograph, we employ a particular saturation plane transformation utilizing the function $f(x)$. This method entails breaking down the image into its hue, saturation, and value components. The transformation is then applied to the saturation plane, with the parameter 'a' adjusted to achieve the best outcomes. Following this adjustment, the modified components are combined back together, resulting in an enhanced version of the original image with increased vibrance. To provide a comprehensive visual comparison, we display both the enhanced image and the transformation graph alongside the original.

$$f(x) = \min \left(x + \frac{a}{128} e^{-\frac{(x-128)^2}{2\sigma^2}}, 255 \right)$$

```

# Load the image in BGR color space
image_bgr = cv.imread('images/spider.png',
cv.IMREAD_COLOR)

# Convert the BGR image to HSV color space
image_hsv = cv.cvtColor(image_bgr, cv.COLOR_BGR2HSV)

# Extract the saturation (S) plane
saturation_plane = image_hsv[:, :, 1]

# Define the intensity transformation function
def intensity_transform(x, a, sigma=70):
    # Apply an intensity transformation to enhance
    vibrance
    transformed_x = np.clip(x + a * 128 * np.exp(-(x -
128) ** 2 / (2 * sigma ** 2)), 0, 255)
    return transformed_x

# Apply the intensity transform function to enhance
saturation
transformed_saturation_plane =
intensity_transform(saturation_plane, 0.5)

# Create a copy of the HSV image for modification
image_copy = image_hsv.copy()

# Recombine the planes with the transformed saturation
image_copy[:, :, 1] = transformed_saturation_plane

# Convert the HSV image back to BGR color space
original_image = cv.cvtColor(image_hsv, cv.COLOR_HSV2RGB)
transformed_image = cv.cvtColor(image_copy,
cv.COLOR_HSV2RGB)

```




figure 8:- Original and Enhanced

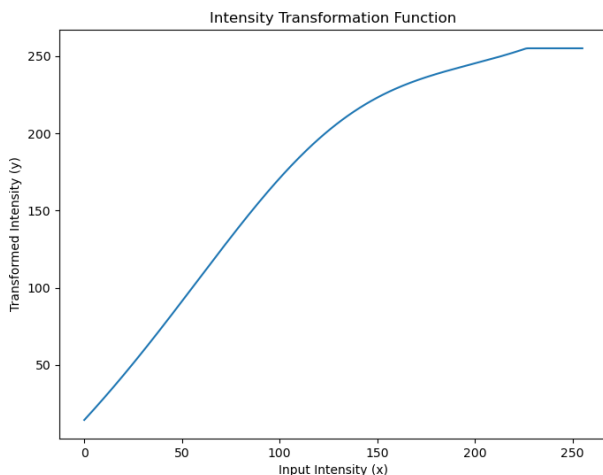


figure 9:- Intensity transformation

5. Question 05

This assignment requires the development of a custom Python function for histogram equalization applied to a given image. This process is designed to exclude the use of built-in functions like 'cv2.equalizeHist()'. The function begins by loading the image as a NumPy array and subsequently computes its histogram and cumulative distribution function (CDF). Following this, the CDF is normalized to ensure it falls within the desired range. The core step involves redistributing pixel intensities according to the CDF. Finally, the histograms before and after equalization are displayed, providing insight into how the transformation impacts the image's intensity distribution.

important code related to histogram equalization:

```

def compute_histogram(image):
    # Compute the histogram of the input image
    hist, bins = np.histogram(image.ravel(), 256, [0, 256])
    return hist, bins

def compute_cumulative_histogram(hist):
    # Compute the cumulative histogram
    cumulative_hist = hist.cumsum()
    normalized_cumulative_hist = cumulative_hist *
hist.max() / cumulative_hist.max()
    return cumulative_hist, normalized_cumulative_hist

def compute_probabilities(hist):
    # Compute probabilities for each intensity level
    probabilities = hist / np.sum(hist)
    return probabilities

def compute_cumulative_sum(hist):
    # Compute the cumulative sum of pixels
    cumulative_sum = np.zeros(256)
    for i in range(len(cumulative_sum)):
        cumulative_sum[i] = np.sum(hist[:i])
    return cumulative_sum

def equalize_histogram(image, cumulative_sum):
    # Equalize the image using the cumulative sum
    equalized_cumulative_sum = np.zeros(256)
    for x in range(len(equalized_cumulative_sum)):
        equalized_cumulative_sum[x] = (cumulative_sum[x]
* 255) / image.size
    equalized_cumulative_sum =
equalized_cumulative_sum.astype('uint8')
    equalized_image = np.zeros(image.shape)
    for i in range(len(image)):

```

```

for j in range(len(image[i])):
    equalized_image[i][j] =
equalized_cumulative_sum[image[i][j]]
equalized_image = equalized_image.astype('uint8')
return equalized_image

```




figure 10:- Original and Equalized




figure 11:- Original and Equalized
Histograms

6. Question 06

This task involves improving the histogram exclusively for the foreground of an image, resulting in an equalized histogram for the image's foreground. The process consists of several steps, which include segmenting the image into its hue, saturation, and value components, creating a foreground mask through thresholding, and using operations such as 'cv2.bitwise_and' to isolate the foreground region. Next, the histogram for the foreground is calculated, and its cumulative sum is determined. The application of histogram equalization formulas enhances histogram for the foreground. Finally, the background information is reintegrated to achieve the desired result. The key components of this task involve the hue, saturation, and value planes, the foreground mask, the original image, and the resulting image displaying the histogram-equalized foreground.

This code implementation of this task.

```

# Convert the image to the HSV color space (Hue,
Saturation, Value)
hsv_image = cv.cvtColor(image, cv.COLOR_BGR2HSV)
# Split the HSV image into its individual channels: Hue,
Saturation, and Value
hue_channel, saturation_channel, value_channel =
cv.split(hsv_image)
# Display the individual channels in grayscale
plt.figure(figsize=(12, 4))
# Manually set a threshold value (adjust as needed)
threshold_value = 14
# Choose the saturation channel for thresholding
threshold_channel = saturation_channel
# Perform thresholding to extract the foreground mask
_, foreground_mask = cv.threshold(threshold_channel,
threshold_value, 255, cv.THRESH_BINARY)
# Extract the foreground region using cv2.bitwise_and
foreground = cv.bitwise_and(image, image,
mask=foreground_mask)
# Compute the histogram of the foreground region
foreground_gray = cv.cvtColor(foreground,
cv.COLOR_BGR2GRAY)
hist, bins = np.histogram(foreground_gray.ravel(),
bins=256, range=[0, 255])
# Calculate the cumulative sum of the histogram
cdf = hist.cumsum()
# Perform histogram equalization on the foreground
equalized_foreground_gray =
cv.equalizeHist(foreground_gray).astype(np.uint8)
# Convert the equalized foreground back to BGR format
equalized_foreground =
cv.cvtColor(equalized_foreground_gray, cv.COLOR_GRAY2BGR)
# Extract the background by inverting the foreground mask
background_mask = cv.bitwise_not(foreground_mask)
background = cv.bitwise_and(image, image,
mask=background_mask)
# Combine the equalized foreground with the background
result = cv.add(equalized_foreground, background)
# Display the original image and the result image side by
side
image_rgb = cv.cvtColor(image, cv.COLOR_BGR2RGB)
result_rgb = cv.cvtColor(result, cv.COLOR_BGR2RGB)

```



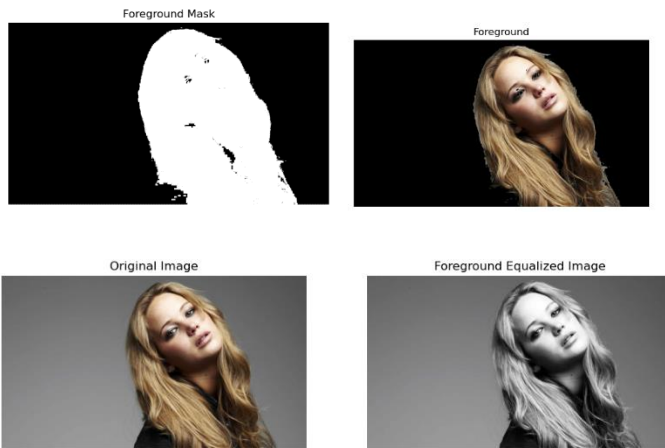


figure 12:- h, s, v planes and Original, Foreground and Equalized Images

7. Question 07

In this section, we explore various techniques for applying Sobel filtering to an image:

- In the initial approach, we employ the existing 'filter2D' function to execute Sobel filtering.
- In the second method, we craft our own custom code for Sobel filtering, granting us greater control over the process.
- Finally, we leverage a convolution property using a specific matrix configuration. By applying this matrix convolution to the image, we achieve the Sobel filtering effect, which effectively enhances and emphasizes the edges within the image.

Here is the code snippet that corresponds to the first approach:(for two different kernels)

```
# Define two convolution kernels for image filtering
ker1=np.array([[ -1,0,1],[-2,0,2],[ -1,0,1]])
ker2=np.array([[ -1,-2,-1],[0,0,0],[1,2,1]])
# Apply the convolution operation with the first kernel
img1 = cv.filter2D(img, -1, ker1)
# Apply the convolution operation with the second kernel
img2 = cv.filter2D(img, -1, ker2)
```




figure 13:- Sobel Filter Using "filter2D" function

Here is the code snippet that corresponds to the second approach:(for two different kernels)

```
# Define Sobel kernels
ker1 = np.array([[ -1, 0, 1], [-2, 0, 2], [-1, 0, 1]])
ker2 = np.array([[ -1, -2, -1], [0, 0, 0], [1, 2, 1]])
# Function to apply convolution filter
def filter(image, kernel):
    assert kernel.shape[0] % 2 == 1 and kernel.shape[1] % 2 == 1
    k_hh, k_hw = math.floor(kernel.shape[0] / 2),
    math.floor(kernel.shape[1] / 2)
    h, w = image.shape
    image_float = cv.normalize(image.astype(float), None,
    0.0, 1.0, cv.NORM_MINMAX)
    result = np.zeros(image.shape, float)
    for m in range(k_hh, h - k_hh):
        for n in range(k_hw, w - k_hw):
            result[m, n] = np.dot(image_float[m - k_hh:m
            + k_hh + 1, n - k_hw:n + k_hw + 1].flatten(),
            kernel.flatten())
    return result
# Apply Sobel filters
img_1 = filter(img, ker1)
img_2 = filter(img, ker2)
```




figure 14:- Sobel filtering without using "filter2D" function.

Here is the code snippet that corresponds to the third approach:

```
kernel1 = np.array([[1], [2], [1]])
kernel2 = np.array([[1, 0, -1]])
kernel = kernel1*kernel2
# Apply the 1D convolutions
conv1 = cv.filter2D(img, -1, kernel)
```

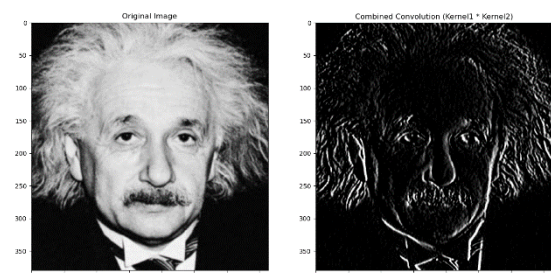


figure 15:- Sobel filtering using third approach.

8. Question 08

This code is designed for image magnification with a specified factor 's' within the (0, 10) range. It offers two distinct zooming techniques: nearest-neighbor and bilinear interpolation. When utilizing the nearest-neighbor method, the program enlarges the image by replicating the nearest pixel values, preserving a blocky appearance. In contrast, the bilinear interpolation method calculates pixel values by considering the weighted average of neighboring pixels, resulting in smoother enlarged images. To assess the accuracy of the zooming implementation, it calculates the normalized sum of squared differences (SSD) when scaling up small images by a factor of 4 and compares them to the original images.

Important code for nearest-neighbor method

```
def zoom_nearest_neighbor(image, factor):
    h, w, _ = image.shape
    new_h = int(h * factor)
    new_w = int(w * factor)
    zoomed_image = np.zeros((new_h, new_w, 3),
dtype=np.uint8)
    for i in range(new_h):
        for j in range(new_w):
            orig_i = int(i / factor)
            orig_j = int(j / factor)
            zoomed_image[i, j] = image[orig_i, orig_j]
    return zoomed_image
```

Important code for Bilinear method

```
def zoom_bilinear(image, factor):
    h, w, _ = image.shape
    new_h = int(h * factor)
    new_w = int(w * factor)
    zoomed_image = np.zeros((new_h, new_w, 3),
dtype=np.uint8)
    for i in range(new_h):
        for j in range(new_w):
            orig_i = i / factor
            orig_j = j / factor
            i1, i2 = int(np.floor(orig_i)),
int(np.ceil(orig_i))
```

```
            j1, j2 = int(np.floor(orig_j)),
int(np.ceil(orig_j))
            i1 = max(0, min(i1, h - 1)) # Ensure indices
stay within image boundaries
            i2 = max(0, min(i2, h - 1))
            j1 = max(0, min(j1, w - 1))
            j2 = max(0, min(j2, w - 1))
            # Bilinear interpolation
            value = (1 - (orig_i - i1)) * (1 - (orig_j -
j1)) * image[i1, j1] + \
                    (1 - (orig_i - i1)) * (orig_j - j1) *
image[i1, j2] + \
                    (orig_i - i1) * (1 - (orig_j - j1)) *
image[i2, j1] + \
                    (orig_i - i1) * (orig_j - j1) *
image[i2, j2]
            zoomed_image[i, j] = value.astype(np.uint8)
    return zoomed_image
```

Important code for Calculate SSD

```
def compute_normalized_ssd(image1, image2):
    pixel_difference = image1 - image2
    squared_pixel_difference = pixel_difference ** 2
    sum_of_squared_differences =
np.sum(squared_pixel_difference)
    normalized_ssd = sum_of_squared_differences /
(image1.size * 255 ** 2)
    return normalized_ssd
```

Image	Nearest_Neighbor	Bilinear
01	0.000481	0.000603
02	0.000183	0.000249
03	-	-
04	0.001210	0.001255
05	0.000777	0.000825
06	0.000469	0.000546
07	0.000430	0.000464
08	-	-
09	0.000325	0.000410
10	-	-
11	-	-

9. Question 09

In the provided image of a flower, where both the foreground and background are in sharp focus, the task comprises two main objectives:

1. In the first part, the "grab-Cut" algorithm is utilized for image segmentation. This results in two

primary outcomes: the creation of a segmentation mask that distinguishes the foreground from the background and the isolation of the foreground and background images.

2. In the second part, the objective is to enhance the image by applying significant background blur. This action aims to create a visually striking separation between the subject and its background. By presenting both the original and enhanced images side by side, viewers can clearly observe the improvements made.

Important code for this task:

```
# Create a mask initialized with zeros
mask = np.zeros(img.shape[:2], np.uint8)
# Define a rectangle that contains the object of interest
(flower)
rectangle = (50, 50, 505, 505) # (x, y, width, height)
# Apply GrabCut algorithm with rectangle initialization
background_model = np.zeros((1, 65), np.float64)
foreground_model = np.zeros((1, 65), np.float64)
cv.grabCut(img, mask, rectangle, background_model,
foreground_model, 5, cv.GC_INIT_WITH_RECT)
# Create a binary mask where the probable background and
definite background are set to 0
mask1 = np.where((mask == 2) | (mask == 0), 0,
1).astype('uint8')
foreground_img = img * mask1[:, :, np.newaxis]
# Apply GrabCut algorithm with mask initialization
cv.grabCut(img, mask, rectangle, background_model,
foreground_model, 5, cv.GC_INIT_WITH_MASK)
# Create a binary mask where the probable foreground and
definite foreground are set to 0
mask2 = np.where((mask == 3) | (mask == 1), 0,
1).astype('uint8')
background_img = img * mask2[:, :, np.newaxis]
```



figure 16:- 6 Segmented, Foreground and Background images

```
# Create a mask initialized with zeros
mask = np.zeros(image.shape[:2], np.uint8)
# Define a rectangle that contains the object of interest
(flower)
```

```
rect = (50, 50, image.shape[1] - 50, image.shape[0] - 50)
# Apply GrabCut algorithm with rectangle initialization
bgdModel = np.zeros((1, 65), np.float64)
fgdModel = np.zeros((1, 65), np.float64)
cv.grabCut(image, mask, rect, bgdModel, fgdModel, 5,
cv.GC_INIT_WITH_RECT)
# Modify the mask to get a binary mask
mask2 = np.where((mask == 2) | (mask == 0), 0,
1).astype('uint8')
mask3 = 1 - mask2
# Apply the mask to get segmented images
foreground_img = image * mask2[:, :, np.newaxis]
background_img = image * mask3[:, :, np.newaxis]
# Create a blurred background
blurred_background = cv.GaussianBlur(background_img, (0,
0), sigmaX=10, sigmaY=10)
blurred_background = blurred_background * mask3[:, :,
np.newaxis]
# Replace the background in the segmented image
enhanced_img = foreground_img + blurred_background
```



figure 17:- Original and Enhanced images

Regarding Part C:

The darkening of the background near the flower's edge in an enhanced image is due to intentional blurring, a technique used to separate the subject from the background. This blurring, often used for creating a bokeh effect, reduces sharpness in distant elements, making the background intentionally blurry and darker. This enhances the flower's prominence as the focal point and creates a visually pleasing contrast between the subject and its surroundings.
